

```
In [21]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

sns.set(style="whitegrid")
```

```
In [22]: file_path = '/content/MathE dataset (4).csv'
dataset = pd.read_csv(file_path, delimiter=';', encoding='ISO-8859-1')
```

use delimiter=';' because of this csv file include ; instead of ,

use encoding='ISO-8859-1' because of this csv file include some characters couldn't be interpreted correctly in UTF-8 format

Data visualisation

Visualize the distribution of students by country

Using a map visualization technique

```
In [23]: import plotly.express as px

country_counts = dataset['Student Country'].value_counts().reset_index()
country_counts.columns = ['Country', 'Count']

# Create a choropleth map
fig = px.choropleth(country_counts,
                    locations="Country",
                    locationmode='country names',
                    color="Count",
                    hover_name="Country",
                    color_continuous_scale=px.colors.sequential.Plasma,
                    title="Student Distribution by Country")

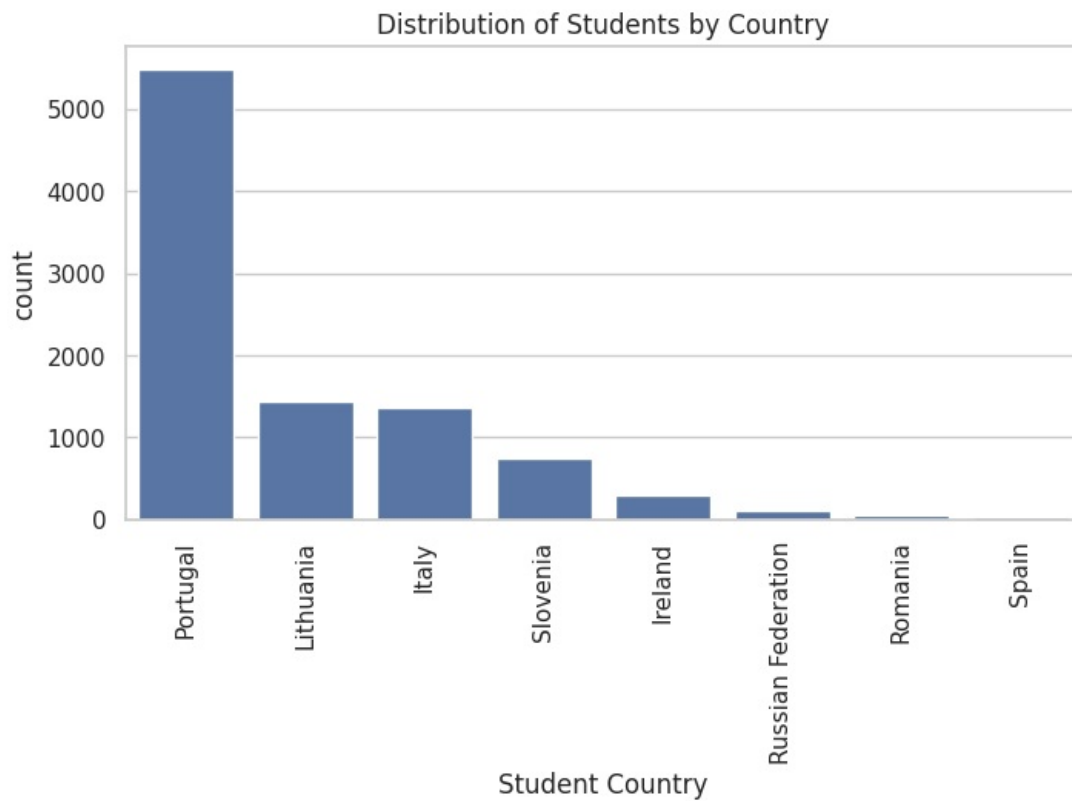
fig.show()
```

this use when dealing with geographical data or location-based analysis

Using chart visualization technique


```
In [24]: plt.figure(figsize=(8, 4))
sns.countplot(data=dataset, x='Student Country', order=dataset['Student Country'].value_counts().index)
plt.title('Distribution of Students by Country')
```

```
plt.xticks(rotation=90)  
plt.show()
```




Visualize the distribution of correct (1) vs incorrect (0) answers(bar plot)

```
In [25]: plt.figure(figsize=(8, 6))  
sns.countplot(data=dataset, x='Type of Answer')  
plt.title('Distribution of Correct vs Incorrect Answers')  
plt.show()
```



Visualize the Performance comparison across Basic and Advanced levels


```
In [26]: plt.figure(figsize=(8, 6))
sns.countplot(data=dataset, x='Question Level', hue='Type of Answer')
plt.title('Performance in Basic vs Advanced Questions')
plt.show()
```



Distribution Analysis

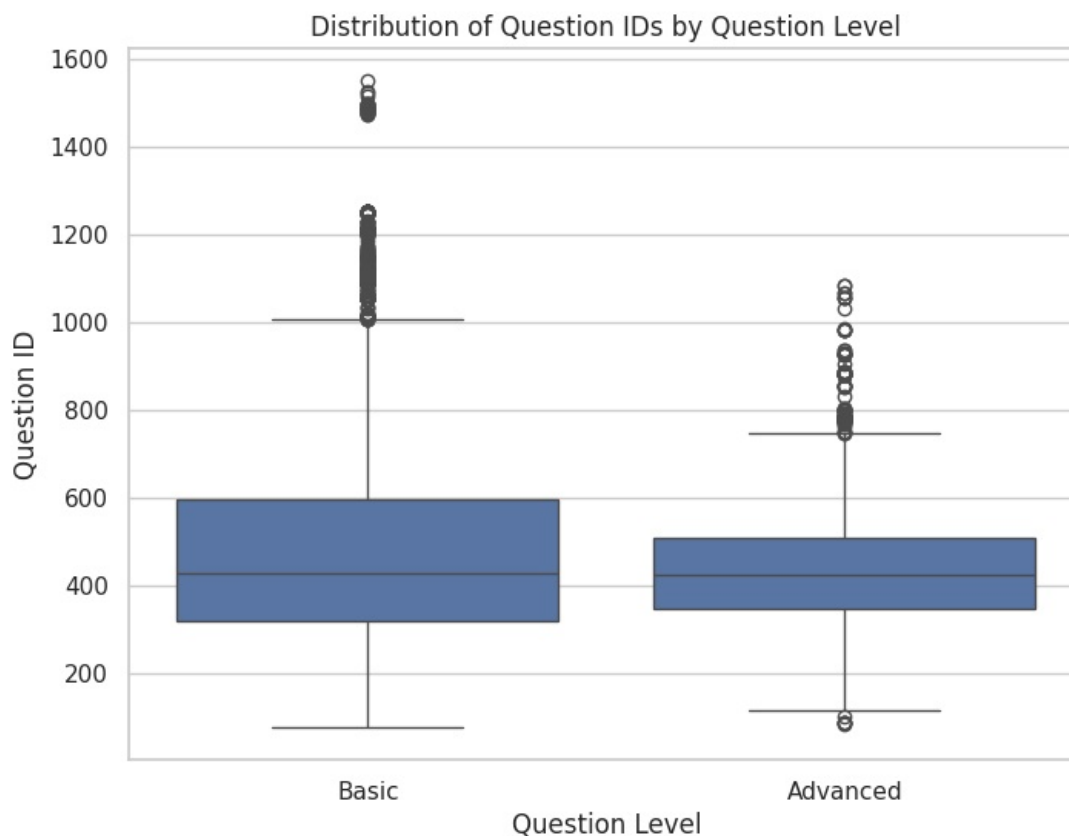
Distribution of Question IDs

```
In [27]: plt.figure(figsize=(8, 6))
plt.hist(dataset['Question ID'], bins=30, color='blue', edgecolor='black')
plt.title('Distribution of Question IDs')
plt.xlabel('Question ID')
plt.ylabel('Frequency')
plt.show()
```



Distribution of Question IDs by Question Level


```
In [28]: plt.figure(figsize=(8, 6))
sns.boxplot(data=dataset, x='Question Level', y='Question ID')
plt.title('Distribution of Question IDs by Question Level')
plt.show()
```



Distribution of questions across topics

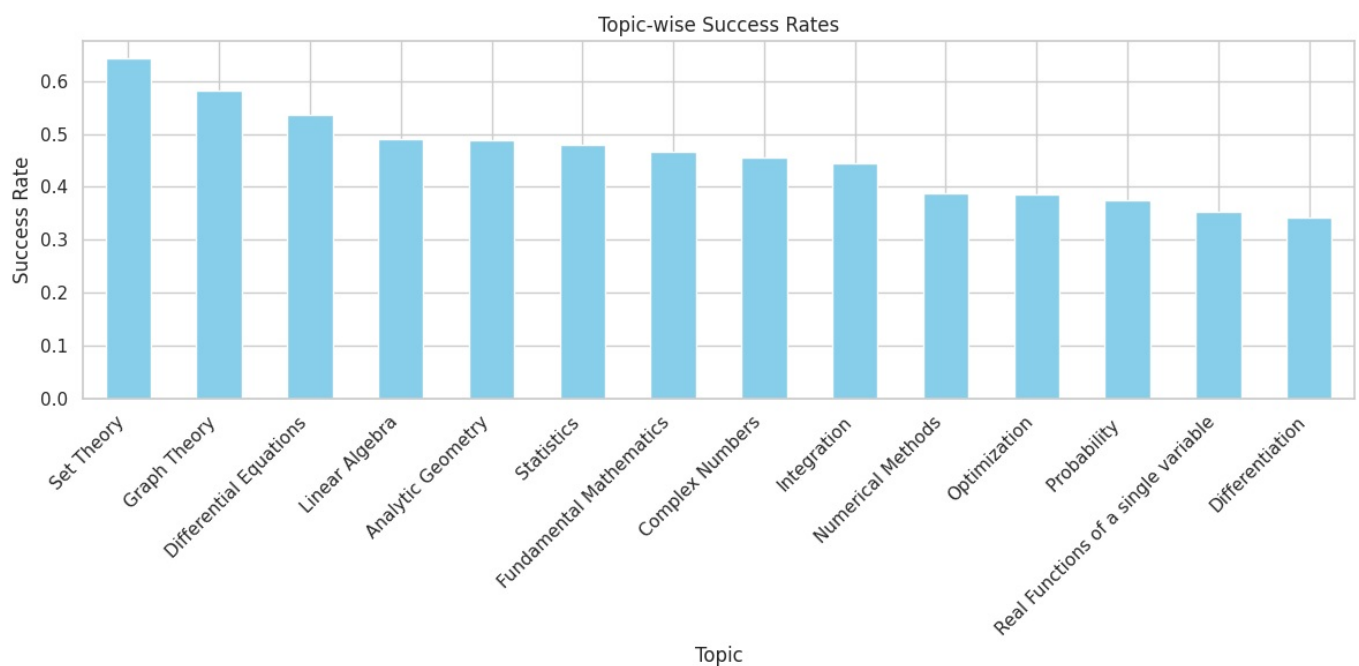
```
In [29]: plt.figure(figsize=(12, 6))
sns.countplot(data=dataset, x='Topic', order=dataset['Topic'].value_counts().index)
plt.title('Distribution of Question Topics')
```

```
plt.xticks(rotation=90)
plt.show()
```




```
In [30]: topic_performance = dataset.groupby('Topic')['Type of Answer'].mean().sort_values(ascending=False)
```

```
In [31]: plt.figure(figsize=(12, 6))
topic_performance.plot(kind='bar', color='skyblue')
plt.title('Topic-wise Success Rates')
plt.xlabel('Topic')
plt.ylabel('Success Rate')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



Distribution of questions across subtopics

```
In [32]: plt.figure(figsize=(12, 6))
sns.countplot(data=dataset, x='Subtopic', order=dataset['Subtopic'].value_counts().index)
plt.title('Distribution of Question Subtopics')
plt.xticks(rotation=90)
plt.show()
```

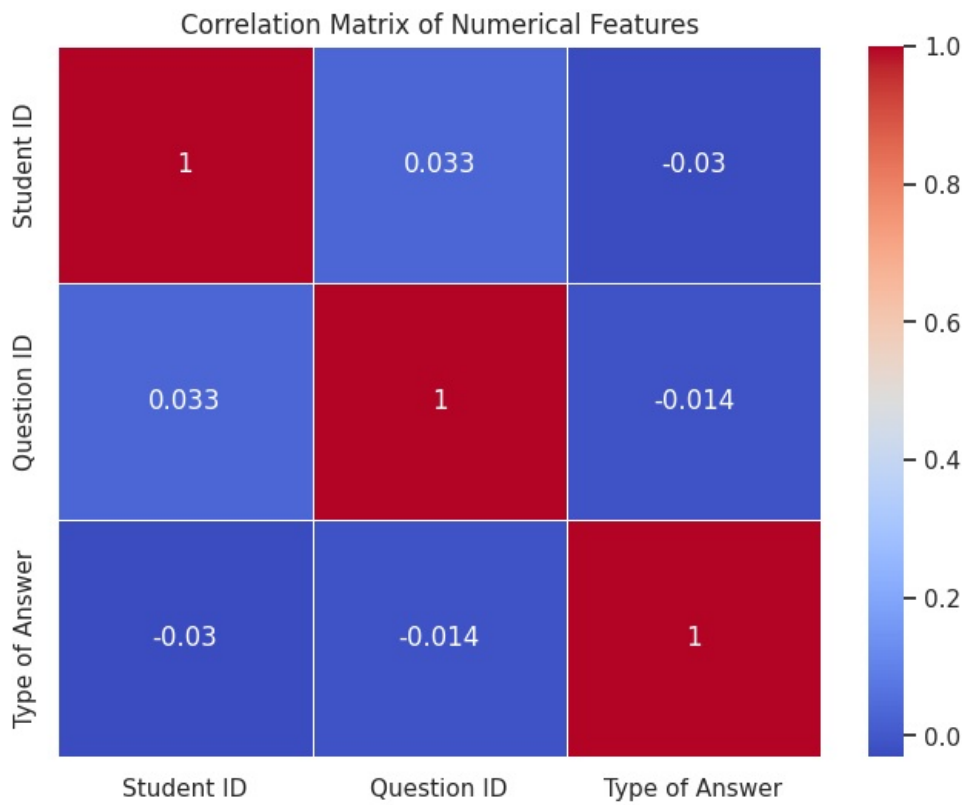


Correlation Analysis

Correlation Heatmap for Numerical Fields

```
In [33]: # Select only numeric columns
numeric_data = dataset.select_dtypes(include=['number'])

plt.figure(figsize=(8, 6))
sns.heatmap(numeric_data.corr(), annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Matrix of Numerical Features')
plt.show()
```

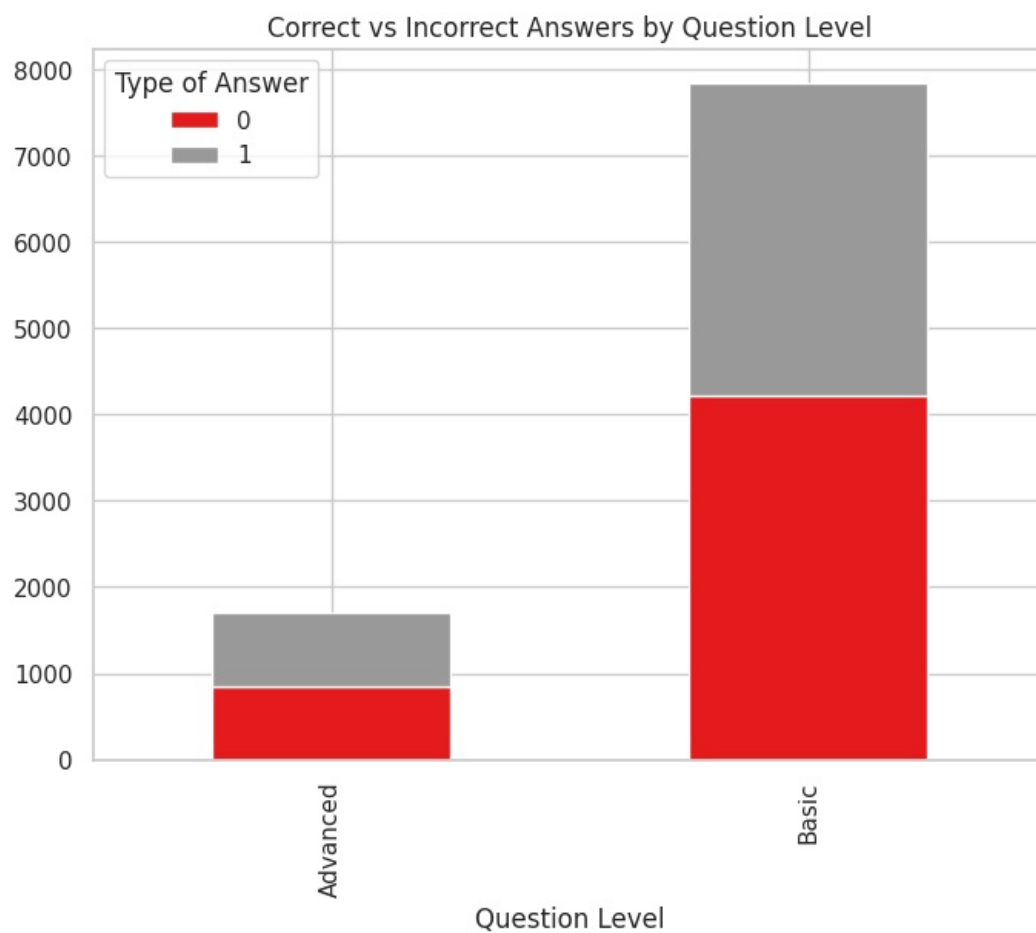


Correlation between Question Level and Answer Type(cross tab)

```
In [34]: cross_tab = pd.crosstab(dataset['Question Level'], dataset['Type of Answer'])
print(cross_tab)

cross_tab.plot(kind='bar', stacked=True, figsize=(8, 6), colormap='Set1')
plt.title('Correct vs Incorrect Answers by Question Level')
plt.show()
```

Type of Answer	0	1
Question Level		
Advanced	849	853
Basic	4227	3617

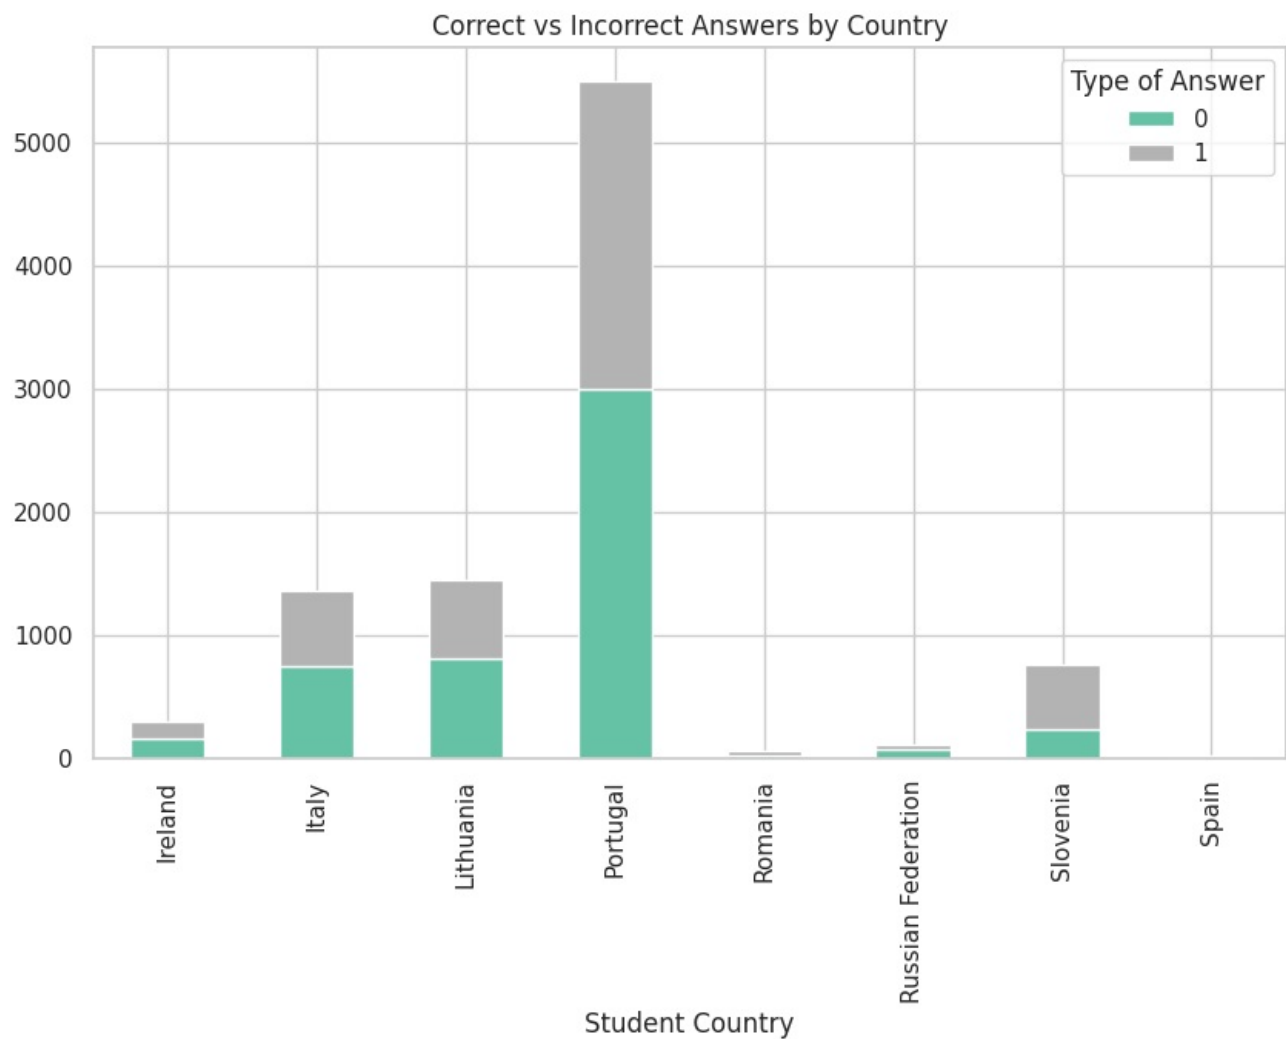


Correlation between Student Country and Answer Type(cross tab)

```
In [35]: country_performance = pd.crosstab(dataset['Student Country'], dataset['Type of Answer'])
print(country_performance)

country_performance.plot(kind='bar', stacked=True, figsize=(10, 6), colormap='Set2')
plt.title('Correct vs Incorrect Answers by Country')
plt.xticks(rotation=90)
plt.show()
```

Type of Answer	0	1
Student Country		
Ireland	162	138
Italy	752	606
Lithuania	814	629
Portugal	3001	2494
Romania	25	35
Russian Federation	70	37
Slovenia	236	519
Spain	16	12



Descriptive statistics of the dataset

```
In [36]: print(dataset.describe(include='all'))
```


	Student ID	Student Country	Question ID	Type of Answer \
count	9546.000000	9546	9546.000000	9546.000000
unique	NaN	8	NaN	NaN
top	NaN	Portugal	NaN	NaN
freq	NaN	5495	NaN	NaN
mean	775.402263	NaN	478.912319	0.468259
std	460.590559	NaN	249.244061	0.499018
min	26.000000	NaN	77.000000	0.000000
25%	380.000000	NaN	323.000000	0.000000
50%	885.000000	NaN	428.000000	0.000000
75%	1219.000000	NaN	571.000000	1.000000
max	1565.000000	NaN	1549.000000	1.000000

	Question Level	Topic	Subtopic \
count	9546	9546	9546
unique	2	14	24
top	Basic	Linear Algebra	Vector Spaces
freq	7844	5726	2749
mean	NaN	NaN	NaN
std	NaN	NaN	NaN
min	NaN	NaN	NaN
25%	NaN	NaN	NaN
50%	NaN	NaN	NaN
75%	NaN	NaN	NaN
max	NaN	NaN	NaN

	Keywords
count	9546
unique	365
top	Linear application,Linearity
freq	443
mean	NaN
std	NaN
min	NaN
25%	NaN
50%	NaN
75%	NaN
max	NaN

Data Preprocessing for Clustering Analysis of Mathematics Learning Patterns

Handling Missing Values

```
In [37]: dataset.isnull().sum()
```

```
Out[37]:
```

	0
Student ID	0
Student Country	0
Question ID	0
Type of Answer	0
Question Level	0
Topic	0
Subtopic	0
Keywords	0

dtype: int64

```
In [38]: from sklearn.preprocessing import LabelEncoder, OneHotEncoder

# 1. Label Encoding
# Create a LabelEncoder object
label_encoder = LabelEncoder()

# Apply label encoding to specified columns
for column in ['Student Country', 'Question Level', 'Topic', 'Subtopic']:
    dataset[column] = label_encoder.fit_transform(dataset[column])

print("Label Encoded Data:")
print(dataset)
```

Label Encoded Data:

	Student ID	Student Country	Question ID	Type of Answer	\
0	647	0	77	0	
1	41	3	77	1	
2	340	3	77	1	
3	641	1	77	0	
4	669	3	77	1	
...	
9541	175	1	1497	1	
9542	175	1	1514	0	
9543	175	1	1521	0	
9544	175	1	1526	1	
9545	175	1	1549	0	

	Question Level	Topic	Subtopic	\
0	1	13	22	
1	1	13	22	
2	1	13	22	
3	1	13	22	
4	1	13	22	
...	
9541	1	1	2	
9542	1	1	2	
9543	1	4	0	
9544	1	4	0	
9545	1	4	0	

Keywords

0	Stem and Leaf diagram,Relative frequency,Sampl...
1	Stem and Leaf diagram,Relative frequency,Sampl...
2	Stem and Leaf diagram,Relative frequency,Sampl...
3	Stem and Leaf diagram,Relative frequency,Sampl...
4	Stem and Leaf diagram,Relative frequency,Sampl...
...	...
9541	Imaginary part,Conjugate number,Modulus of a c...
9542	Operations with complex numbers
9543	Quadratic equations,Simplify expressions,Linea...
9544	Linear equations,Quadratic equations,Simplify ...
9545	Simplify expressions,Linear equations,Quadrati...


[9546 rows x 8 columns]

```
In [39]: from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

```
#scaler = StandardScaler() # For standardization
scaler = MinMaxScaler() # For normalization

# Scale numerical features
numerical_features = ['Question ID']
dataset[numerical_features] = scaler.fit_transform(dataset[numerical_features])
```

```
In [40]: # Plotting the normalized 'Question ID'
plt.figure(figsize=(10, 6))
plt.plot(dataset.index, dataset['Question ID'], marker='o', linestyle='-', color='b', label='Normalized Question ID')
plt.xlabel("Index")
plt.ylabel("Normalized Value")
plt.title("Normalized 'Question ID' Feature")
plt.legend()
plt.grid(True)
plt.show()
```



```
In [41]: print(dataset)
```

	Student ID	Student Country	Question ID	Type of Answer	\
0	647	0	0.000000	0	
1	41	3	0.000000	1	
2	340	3	0.000000	1	
3	641	1	0.000000	0	
4	669	3	0.000000	1	
...	
9541	175	1	0.964674	1	
9542	175	1	0.976223	0	
9543	175	1	0.980978	0	
9544	175	1	0.984375	1	
9545	175	1	1.000000	0	

	Question Level	Topic	Subtopic	\
0	1	13	22	
1	1	13	22	
2	1	13	22	
3	1	13	22	
4	1	13	22	
...	
9541	1	1	2	
9542	1	1	2	
9543	1	4	0	
9544	1	4	0	
9545	1	4	0	

	Keywords
0	Stem and Leaf diagram,Relative frequency,Sampl...
1	Stem and Leaf diagram,Relative frequency,Sampl...
2	Stem and Leaf diagram,Relative frequency,Sampl...
3	Stem and Leaf diagram,Relative frequency,Sampl...
4	Stem and Leaf diagram,Relative frequency,Sampl...
...	...
9541	Imaginary part,Conjugate number,Modulus of a c...
9542	Operations with complex numbers
9543	Quadratic equations,Simplify expressions,Linea...
9544	Linear equations,Quadratic equations,Simplify ...
9545	Simplify expressions,Linear equations,Quadрати...

[9546 rows x 8 columns]

```
In [42]: print(dataset.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9546 entries, 0 to 9545
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Student ID            9546 non-null   int64
1   Student Country       9546 non-null   int64
2   Question ID           9546 non-null   float64
3   Type of Answer        9546 non-null   int64
4   Question Level        9546 non-null   int64
5   Topic                 9546 non-null   int64
6   Subtopic              9546 non-null   int64
7   Keywords              9546 non-null   object
dtypes: float64(1), int64(6), object(1)
memory usage: 596.8+ KB
None
```

```
In [43]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from scipy import stats

# Calculate student performance metrics
student_performance = dataset.groupby('Student ID')['Type of Answer'].agg([
    'mean',      # success rate
    'count'      # number of attempts
]).reset_index()

# Calculate Z-scores for outlier detection
z_scores = stats.zscore(student_performance[['mean', 'count']])

# Define outliers (|z| > 3)
outliers_mask = (abs(z_scores) > 3).any(axis=1)
outliers = student_performance[outliers_mask]
normal = student_performance[~outliers_mask]

# Create the visualization
plt.figure(figsize=(10, 6))


# Plot normal points
plt.scatter(normal['count'],
            normal['mean'],
            alpha=0.5,
            color='blue',
            label='Normal')

# Plot outlier points
plt.scatter(outliers['count'],
            outliers['mean'],
            color='red',
            alpha=0.7,
            label='Outliers')

# Add labels and title
plt.xlabel('Number of Attempts')
plt.ylabel('Success Rate')
plt.title('Student Performance Outliers Analysis')
plt.legend()

# Show the plot
plt.show()

# Print summary statistics
print("\nOutlier Summary:")
print(f"Number of outliers detected: {len(outliers)}")
print(f"Percentage of outliers: {(len(outliers)/len(student_performance))*100:.2f}%")
print("\nOutlier Statistics:")
print(outliers.describe())
```



Outlier Summary:
 Number of outliers detected: 3
 Percentage of outliers: 0.81%

Outlier Statistics:

	Student ID	mean	count
count	3.000000	3.000000	3.000000
mean	493.000000	0.574412	423.000000
std	630.776506	0.150739	267.977611
min	91.000000	0.402715	221.000000
25%	129.500000	0.519114	271.000000
50%	168.000000	0.635514	321.000000
75%	694.000000	0.660260	524.000000
max	1220.000000	0.685007	727.000000

```
In [44]: # Calculate student performance metrics
student_performance = dataset.groupby('Student ID')['Type of Answer'].agg([
    'mean',      # success rate
    'count'      # number of attempts
]).reset_index()

# Calculate Z-scores for outlier detection
z_scores = stats.zscore(student_performance[['mean', 'count']])

# Define outliers (|z| > 3)
outliers_mask = (abs(z_scores) > 3).any(axis=1)

# Get list of Student IDs that are not outliers
normal_student_ids = student_performance[~outliers_mask]['Student ID']

# Create cleaned dataset by filtering out outlier students
cleaned_dataset = dataset[dataset['Student ID'].isin(normal_student_ids)]

# Print summary of data cleaning
print("Original dataset shape:", dataset.shape)
print("Cleaned dataset shape:", cleaned_dataset.shape)
print("Number of records removed:", len(dataset) - len(cleaned_dataset))
print("Percentage of data removed: {:.2f}%".format(
    (len(dataset) - len(cleaned_dataset)) / len(dataset) * 100
))

# Visualize the impact of outlier removal
plt.figure(figsize=(12, 5))

# Before cleaning
plt.subplot(1, 2, 1)
sns.boxplot(data=dataset, y='Type of Answer', x='Question Level')
plt.title('Before Outlier Removal')

# After cleaning
plt.subplot(1, 2, 2)
```

```

sns.boxplot(data=cleaned_dataset, y='Type of Answer', x='Question Level')
plt.title('After Outlier Removal')

plt.tight_layout()
plt.show()

# Save cleaned dataset if needed
# cleaned_dataset.to_csv('cleaned_MathE_dataset.csv', index=False)

# Display summary statistics before and after
print("\nSummary Statistics Comparison:")
print("\nBefore cleaning:")
print(dataset['Type of Answer'].describe())
print("\nAfter cleaning:")
print(cleaned_dataset['Type of Answer'].describe())

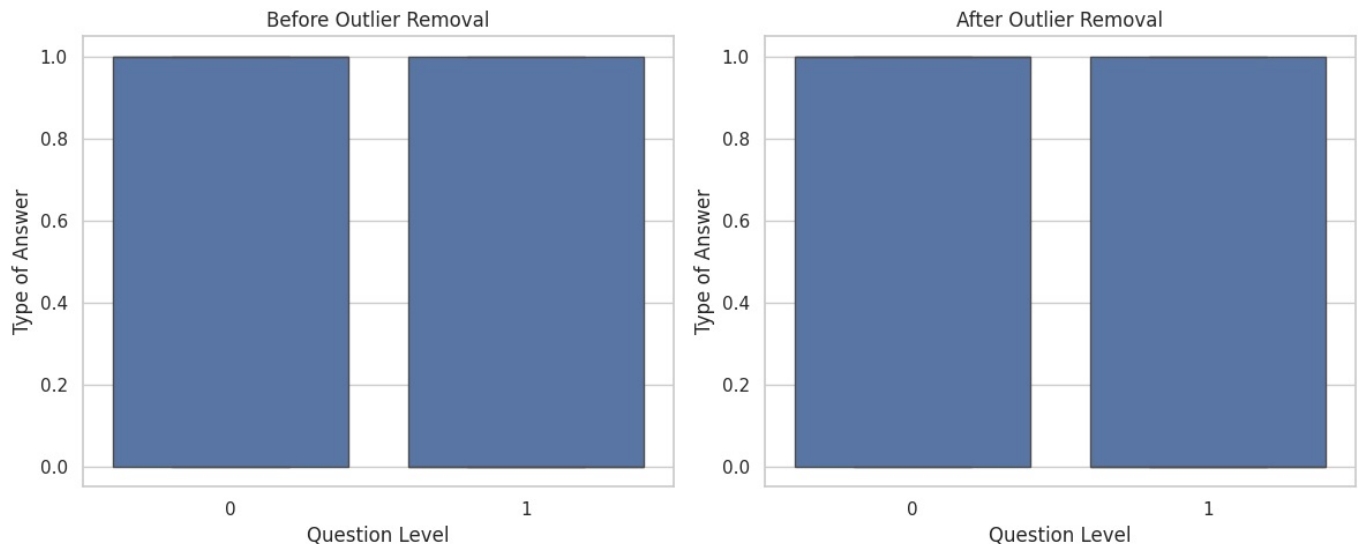
```

Original dataset shape: (9546, 8)

Cleaned dataset shape: (8277, 8)

Number of records removed: 1269

Percentage of data removed: 13.29%



Summary Statistics Comparison:

Before cleaning:

```

count    9546.000000
mean      0.468259
std       0.499018
min       0.000000
25%      0.000000
50%      0.000000
75%      1.000000
max       1.000000

```

Name: Type of Answer, dtype: float64

After cleaning:

```

count    8277.000000
mean      0.444485
std       0.496939
min       0.000000
25%      0.000000
50%      0.000000
75%      1.000000
max       1.000000

```

Name: Type of Answer, dtype: float64

```

In [45]: student_performance = dataset.groupby('Student ID')['Type of Answer'].agg([
        'mean',      # success rate
        'count'      # number of attempts
    ]).reset_index()

# Calculate Z-scores for outlier detection
z_scores = stats.zscore(student_performance[['mean', 'count']])

# Define outliers (|z| > 3)
outliers_mask = (abs(z_scores) > 3).any(axis=1)

# Get list of Student IDs that are not outliers
normal_student_ids = student_performance[~outliers_mask]['Student ID']

# Create cleaned dataset by filtering out outlier students
cleaned_dataset = dataset[dataset['Student ID'].isin(normal_student_ids)]

# Print summary of data cleaning

```

```


print("Original dataset shape:", dataset.shape)
print("Cleaned dataset shape:", cleaned_dataset.shape)
print("Number of records removed:", len(dataset) - len(cleaned_dataset))
print("Percentage of data removed: {:.2f}%".format(
    (len(dataset) - len(cleaned_dataset)) / len(dataset) * 100
))

# Select only numeric columns for correlation analysis
numeric_columns = cleaned_dataset.select_dtypes(include=[np.number]).columns
correlation_matrix = cleaned_dataset[numeric_columns].corr()

# Visualize correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix of Numeric Features (After Outlier Removal)')
plt.show()

```


Original dataset shape: (9546, 8)
 Cleaned dataset shape: (8277, 8)
 Number of records removed: 1269
 Percentage of data removed: 13.29%



```

In [46]: # Performance analysis with cleaned data
plt.figure(figsize=(10, 6))
sns.countplot(data=cleaned_dataset, x='Question Level', hue='Type of Answer')
plt.title('Performance Distribution After Outlier Removal')
plt.show()


```



```
In [47]: # Country-wise performance analysis
country_performance = cleaned_dataset.groupby('Student Country')['Type of Answer'].mean().sort_values(ascending=
```

```
In [48]: # Visualize country-wise performance
plt.figure(figsize=(12, 6))
country_performance.plot(kind='bar')
plt.title('Average Performance by Country (After Outlier Removal)')
plt.xlabel('Country')
plt.ylabel('Average Success Rate')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Print summary statistics
print("\nCountry-wise Performance Summary:")
print(country_performance)
```



Country-wise Performance Summary:
Student Country
6 0.750000
4 0.583333
0 0.460000
3 0.456011
1 0.446244
7 0.428571
2 0.378788
5 0.345794
Name: Type of Answer, dtype: float64

```
In [49]: # Additional analysis: Question Level performance
level_performance = cleaned_dataset.groupby('Question Level')['Type of Answer'].agg([
    'mean',
    'count',
    'std'
]).round(3)

print("\nQuestion Level Performance Summary:")
print(level_performance)
```

Question Level Performance Summary:

	mean	count	std
Question Level			
0	0.503	1530	0.500
1	0.431	6747	0.495

Feature Selection

```
In [50]: import pandas as pd
import numpy as np
from sklearn.feature_selection import mutual_info_classif, RFE
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression, Lasso
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split

# Load cleaned dataset
dataset = cleaned_dataset.copy()

# Define target and features
target = 'Type of Answer' # Target variable
X = dataset.drop(columns=[target, 'Student ID']) # Exclude target and IDs
y = dataset[target]

# Ensure all columns are numeric
for column in X.select_dtypes(include=['object']).columns:
    X[column] = LabelEncoder().fit_transform(X[column])

# Train-test split for evaluation purposes
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Dictionary to store selected features
selected_features = {}

# 1. Filter Method: Correlation Analysis
correlation = X.corrwith(y).abs().sort_values(ascending=False)
selected_features['correlation'] = correlation.index.tolist()[:5] # Top 5 features
print("\nTop features by Correlation:")
print(correlation.head())

# 2. Filter Method: Mutual Information
mi_scores = mutual_info_classif(X, y, random_state=42)
mi_scores_df = pd.Series(mi_scores, index=X.columns).sort_values(ascending=False)
selected_features['mutual_info'] = mi_scores_df.index.tolist()[:5]
print("\nTop features by Mutual Information:")
print(mi_scores_df.head())

# 3. Wrapper Method: Recursive Feature Elimination (RFE)
log_model = LogisticRegression(max_iter=1000, random_state=42)
rfe = RFE(log_model, n_features_to_select=5)
rfe.fit(X_train, y_train)
selected_features['rfe'] = X.columns[rfe.support_].tolist()
print("\nTop features by Recursive Feature Elimination:")
print(selected_features['rfe'])

# 4. Embedded Method: Random Forest Importance
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
feature_importances = pd.Series(rf_model.feature_importances_, index=X.columns).sort_values(ascending=False)
selected_features['random_forest'] = feature_importances.index.tolist()[:5]
```

```

print("\nTop features by Random Forest Importance:")
print(feature_importances.head())

# 5. Embedded Method: Lasso (L1 Regularization)
lasso_model = Lasso(alpha=0.01, max_iter=1000, random_state=42)
lasso_model.fit(X_train, y_train)
lasso_coef = pd.Series(lasso_model.coef_, index=X.columns)
selected_features['lasso'] = lasso_coef[lasso_coef != 0].index.tolist()
print("\nTop features by LASSO Regularization:")
print(selected_features['lasso'])

# Summary of Selected Features
print("\nSummary of Selected Features:")
for method, features in selected_features.items():
    print(f"{method}: {features}")

```

Top features by Correlation:

```

Question Level    0.055707
Student Country   0.016097
Topic             0.010830
Question ID       0.004610
Subtopic          0.001733
dtype: float64

```

Top features by Mutual Information:

```

Question ID       0.020590
Keywords          0.009564
Student Country   0.005757
Topic             0.005664
Question Level    0.001301
dtype: float64

```

Top features by Recursive Feature Elimination:

```
['Student Country', 'Question ID', 'Question Level', 'Topic', 'Subtopic']
```

Top features by Random Forest Importance:

```

Question ID       0.473920
Keywords          0.239359
Student Country   0.167703
Subtopic          0.058453
Topic             0.043578
dtype: float64

```

Top features by LASSO Regularization:

```
['Question Level', 'Topic', 'Subtopic', 'Keywords']
```

Summary of Selected Features:

```

correlation: ['Question Level', 'Student Country', 'Topic', 'Question ID', 'Subtopic']
mutual_info: ['Question ID', 'Keywords', 'Student Country', 'Topic', 'Question Level']
rfe: ['Student Country', 'Question ID', 'Question Level', 'Topic', 'Subtopic']
random_forest: ['Question ID', 'Keywords', 'Student Country', 'Subtopic', 'Topic']
lasso: ['Question Level', 'Topic', 'Subtopic', 'Keywords']

```

```

In [12]: # Summary statistics for selected features
selected_features = ['Question Level', 'Student Country', 'Question ID', 'Topic', 'Subtopic']

print("Descriptive Statistics for Selected Features:")
print(dataset[selected_features + ['Type of Answer']].describe())

```

Descriptive Statistics for Selected Features:

	Question Level	Student Country	Question ID	Topic	Subtopic \
count	8277.000000	8277.000000	8277.000000	8277.000000	8277.000000
mean	0.815150	2.484354	0.278216	6.154766	14.264226
std	0.388199	0.997897	0.179041	2.865115	7.867845
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	2.000000	0.164402	4.000000	8.000000
50%	1.000000	3.000000	0.238451	7.000000	15.000000
75%	1.000000	3.000000	0.377717	7.000000	23.000000
max	1.000000	7.000000	1.000000	13.000000	23.000000

	Type of Answer
count	8277.000000
mean	0.444485
std	0.496939
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

Correlation between selected features and target

```
In [13]: # Correlation between selected features and target
correlations = dataset[selected_features + ['Type of Answer']].corr()
print("\nCorrelation Matrix:")
print(correlations['Type of Answer'].sort_values(ascending=False))
```

Correlation Matrix:

```
Type of Answer    1.000000
Student Country    0.016097
Topic              0.010830
Subtopic          -0.001733
Question ID       -0.004610
Question Level    -0.055707
Name: Type of Answer, dtype: float64
```

Feature Importances from Random Forest

```
In [14]: from sklearn.ensemble import RandomForestClassifier

# Train a Random Forest model
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train[selected_features], y_train)

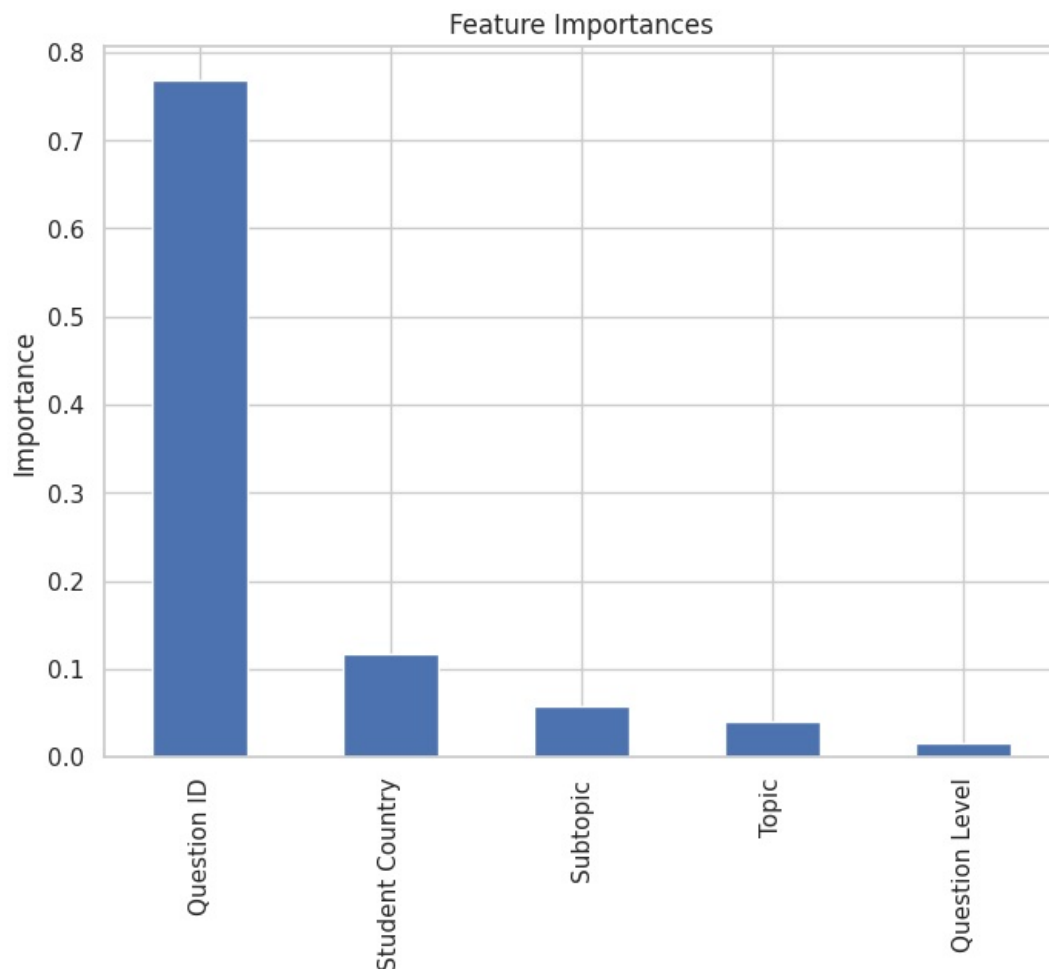
# Feature importance
importances = pd.Series(rf_model.feature_importances_, index=selected_features)
importances.sort_values(ascending=False, inplace=True)

print("\nFeature Importances from Random Forest:")
print(importances)

# Visualization
importances.plot(kind='bar', title='Feature Importances', figsize=(8, 6))
plt.ylabel('Importance')
plt.show()
```

Feature Importances from Random Forest:

```
Question ID    0.768413
Student Country 0.116857
Subtopic       0.057868
Topic          0.041020
Question Level 0.015842
dtype: float64
```



Chi-Square Test

```
In [51]: from scipy.stats import chi2_contingency

print("\nChi-Square Test Results:")
for feature in ['Question Level', 'Student Country', 'Topic', 'Subtopic']:
    contingency_table = pd.crosstab(dataset[feature], dataset['Type of Answer'])
    chi2, p, dof, _ = chi2_contingency(contingency_table)
    print(f"{feature}: Chi2 = {chi2:.3f}, p-value = {p:.3f}")
```

Chi-Square Test Results:

Question Level: Chi2 = 25.398, p-value = 0.000

Student Country: Chi2 = 42.278, p-value = 0.000

Topic: Chi2 = 66.877, p-value = 0.000

Subtopic: Chi2 = 120.045, p-value = 0.000

ANOVA Test for Question ID

```
In [52]: from scipy.stats import f_oneway

anova_result = f_oneway(
    *[dataset['Type of Answer'][dataset['Question ID'] == value]
      for value in dataset['Question ID'].unique()]
)
print(f"\nANOVA Test for Question ID: F-statistic = {anova_result.statistic:.3f}, p-value = {anova_result.pvalue:.3f}")
```

ANOVA Test for Question ID: F-statistic = 1.428, p-value = 0.000

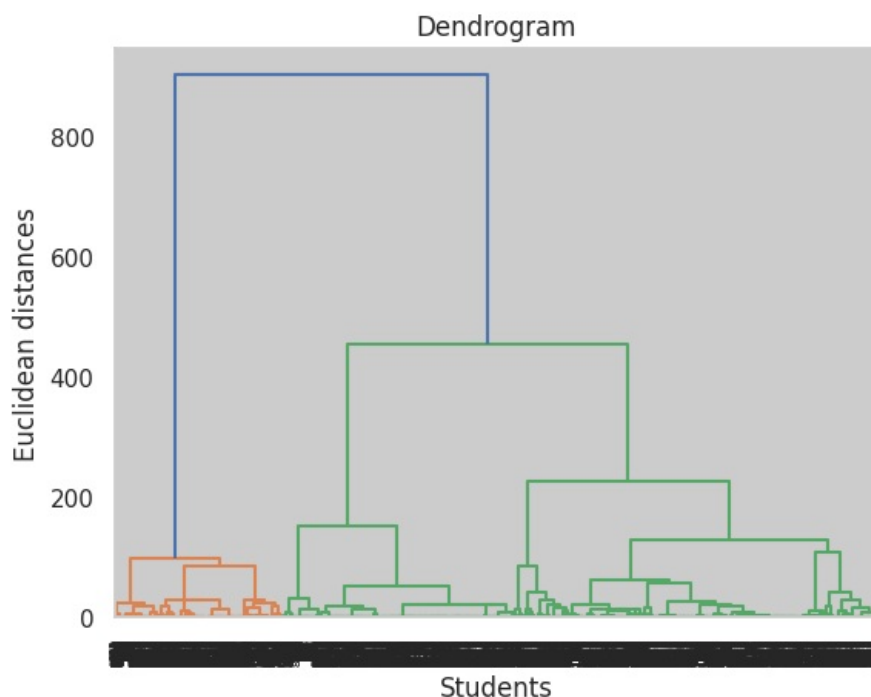
Hierarchical Clustering

```
In [65]: # Hierarchical Clustering

from sklearn.cluster import AgglomerativeClustering
import scipy.cluster.hierarchy as sch

# Select features for clustering
features_for_clustering = ['Question Level', 'Student Country', 'Question ID', 'Topic', 'Subtopic']
X_cluster = cleaned_dataset[features_for_clustering]

# Apply hierarchical clustering
dendrogram = sch.dendrogram(sch.linkage(X_cluster, method='ward'))
plt.title('Dendrogram')
plt.xlabel('Students')
plt.ylabel('Euclidean distances')
plt.show()
```



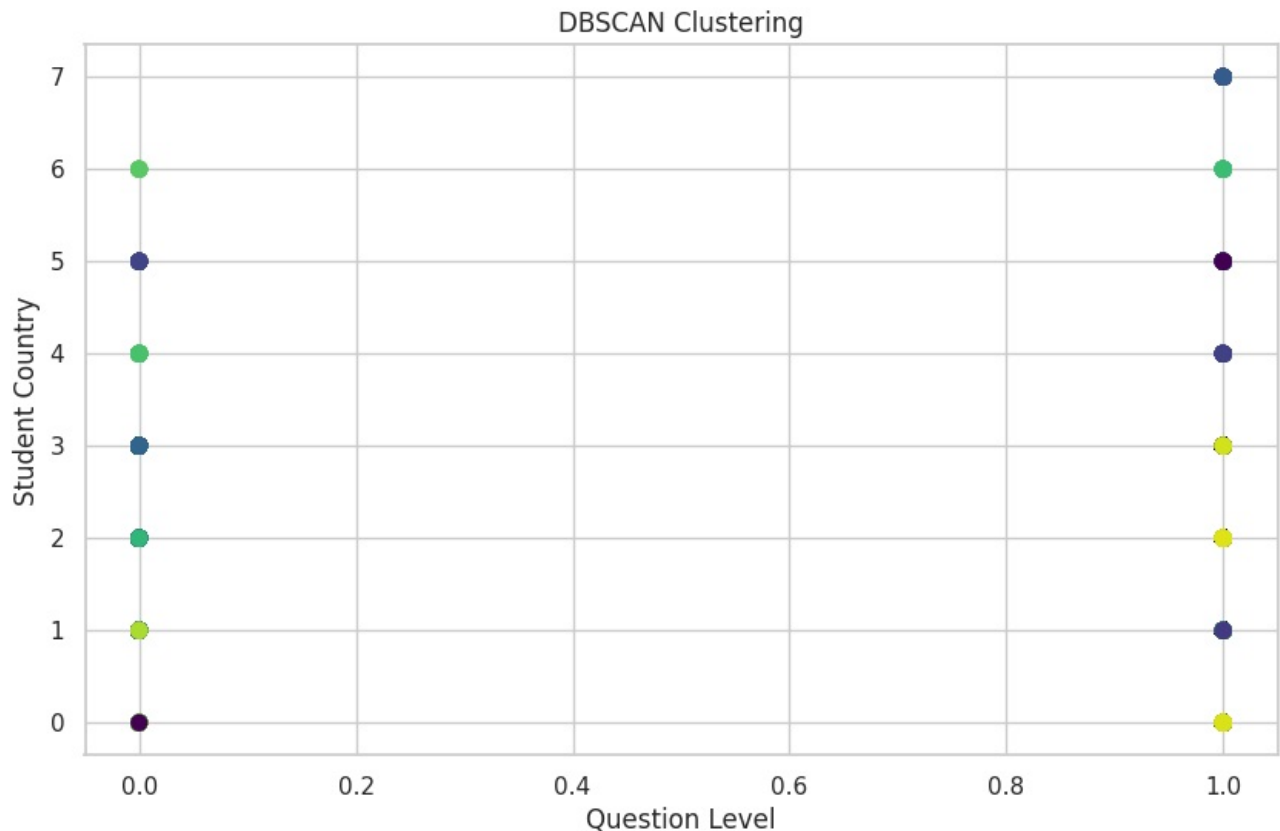
DBSCAN clustering

```
In [16]: from sklearn.cluster import DBSCAN
import pandas as pd
import matplotlib.pyplot as plt

# Select features for clustering
features_for_clustering = ['Question Level', 'Student Country', 'Question ID', 'Topic', 'Subtopic']
X_cluster = cleaned_dataset[features_for_clustering]

# Apply DBSCAN clustering
dbscan = DBSCAN(eps=0.5, min_samples=5)
y_dbscan = dbscan.fit_predict(X_cluster)

# Visualize clusters (with first two features)
plt.figure(figsize=(10, 6))
plt.scatter(X_cluster.iloc[:, 0], X_cluster.iloc[:, 1], c=y_dbscan, s=50, cmap='viridis')
plt.title('DBSCAN Clustering')
plt.xlabel(features_for_clustering[0])
plt.ylabel(features_for_clustering[1])
plt.show()
```



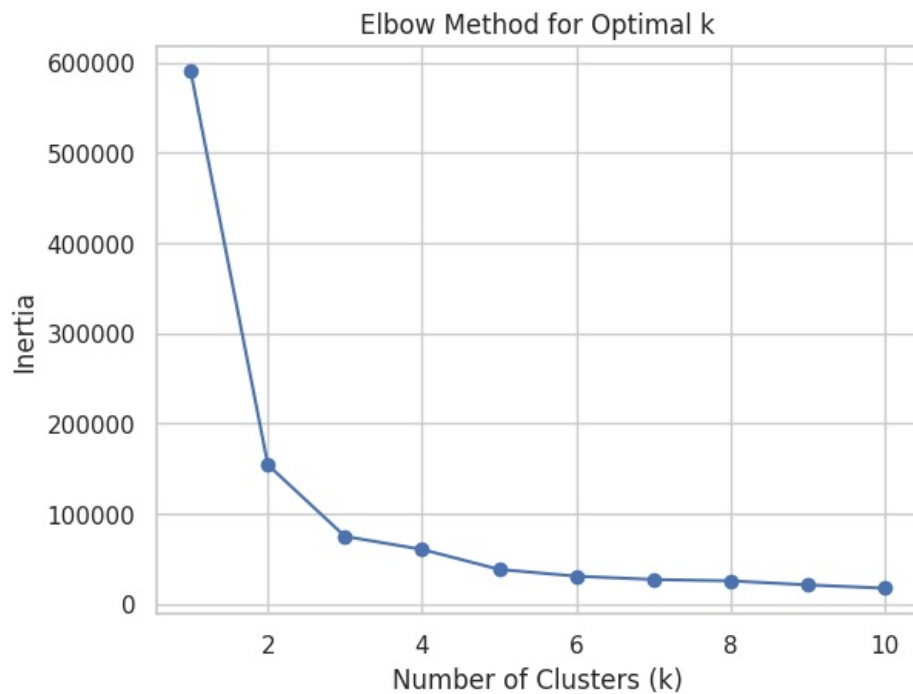
Elbow Method for Optimal k

```
In [62]: from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

# Select features for clustering
cluster_features = ['Question Level', 'Student Country', 'Question ID', 'Topic', 'Subtopic']

# Determine optimal number of clusters (using elbow method)
inertia = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(dataset[cluster_features])
    inertia.append(kmeans.inertia_)

plt.plot(range(1, 11), inertia, marker='o')
plt.title('Elbow Method for Optimal k')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Inertia')
plt.show()
```



```
In [63]: # Perform K-means clustering with optimal k
optimal_k = 3
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
dataset['Cluster'] = kmeans.fit_predict(dataset[cluster_features])
```

```
In [64]: # Evaluate cluster quality (using silhouette score)
silhouette_avg = silhouette_score(dataset[cluster_features], dataset['Cluster'])
print(f"Silhouette Score: {silhouette_avg:.4f}")
```

Silhouette Score: 0.6170

Model Comparison

```
In [59]: # Function to perform and evaluate clustering
def perform_clustering(model, X_cluster, dataset, model_name):
    y_pred = model.fit_predict(X_cluster)
    dataset['Cluster'] = y_pred

    # Visualize clusters (with first two features)
    plt.figure(figsize=(10, 6))
    plt.scatter(X_cluster.iloc[:, 0], X_cluster.iloc[:, 1], c=y_pred, s=50, cmap='viridis')
    plt.title(f'{model_name} Clustering')
    plt.xlabel(features_for_clustering[0])
    plt.ylabel(features_for_clustering[1])
    plt.show()

    # Analyze the clusters
    cluster_analysis = dataset.groupby('Cluster').agg({'Type of Answer': ['mean', 'count']})
    print(f"\n{model_name} Cluster Analysis:\n{cluster_analysis}")

    # Calculate Silhouette Score
    silhouette_avg = silhouette_score(X_cluster, y_pred)
    print(f"{model_name} Silhouette Score: {silhouette_avg:.4f}")

    return dataset

# Select features for clustering
features_for_clustering = ['Question Level', 'Student Country', 'Question ID', 'Topic', 'Subtopic']
X_cluster = cleaned_dataset[features_for_clustering]

# 1. Hierarchical Clustering
hc = AgglomerativeClustering(n_clusters=3, linkage='ward')
cleaned_dataset = perform_clustering(hc, X_cluster, cleaned_dataset.copy(), "Hierarchical")

# 2. DBSCAN Clustering
dbscan = DBSCAN(eps=0.5, min_samples=5) # You might need to tune eps and min_samples
cleaned_dataset = perform_clustering(dbscan, X_cluster, cleaned_dataset.copy(), "DBSCAN")

# 3. K-Means Clustering
```

```
optimal_k = 3 # value from elbow method
```

```
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
```

```
cleaned_dataset = perform_clustering(kmeans, X_cluster, cleaned_dataset.copy(), "KMeans")
```

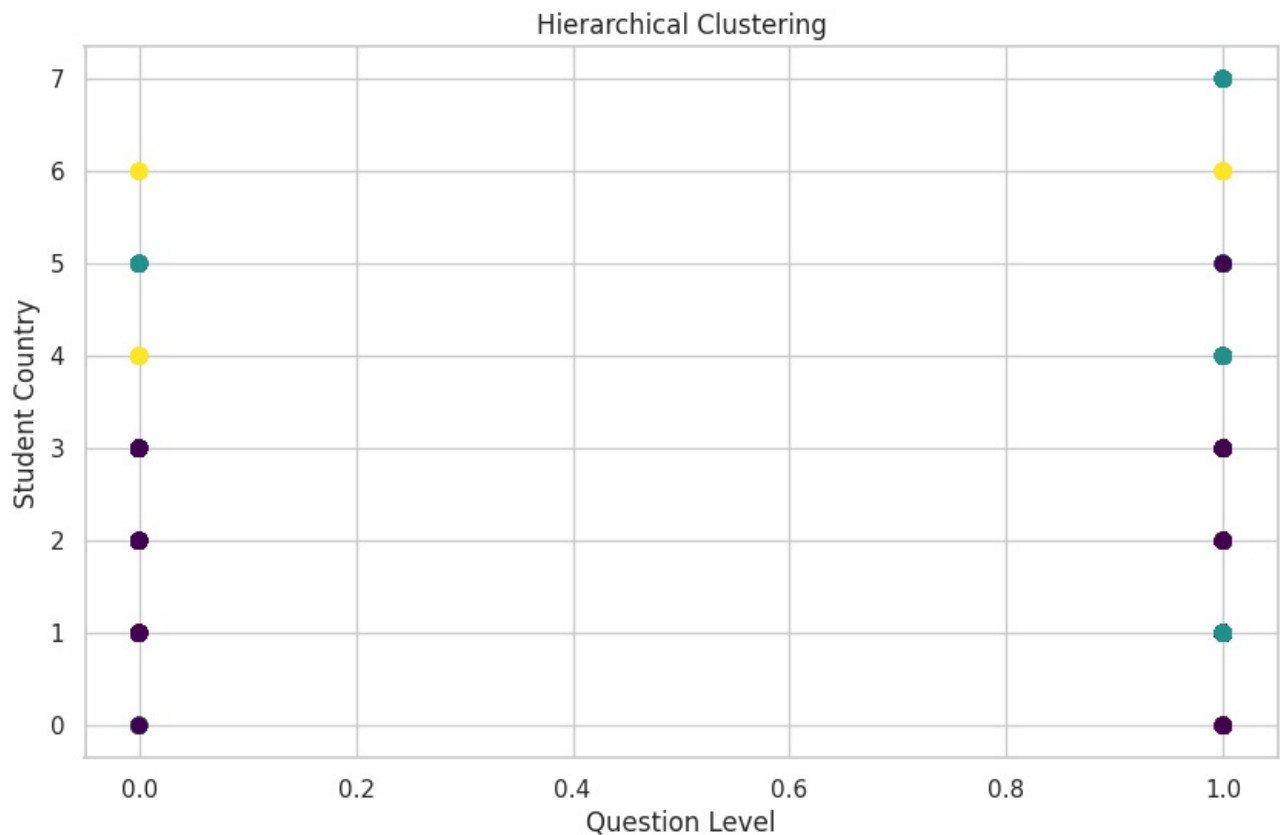
```
# Print a summary of silhouette scores
```

```
print("\nSummary of Silhouette Scores")
```

```
print(f"Hierarchical Clustering: {silhouette_score(X_cluster, hc.fit_predict(X_cluster))}")
```

```
print(f"DBSCAN: {silhouette_score(X_cluster, dbscan.fit_predict(X_cluster))}")
```

```
print(f"KMeans: {silhouette_score(X_cluster, kmeans.fit_predict(X_cluster))}")
```



Hierarchical Cluster Analysis:

Type of Answer

mean count

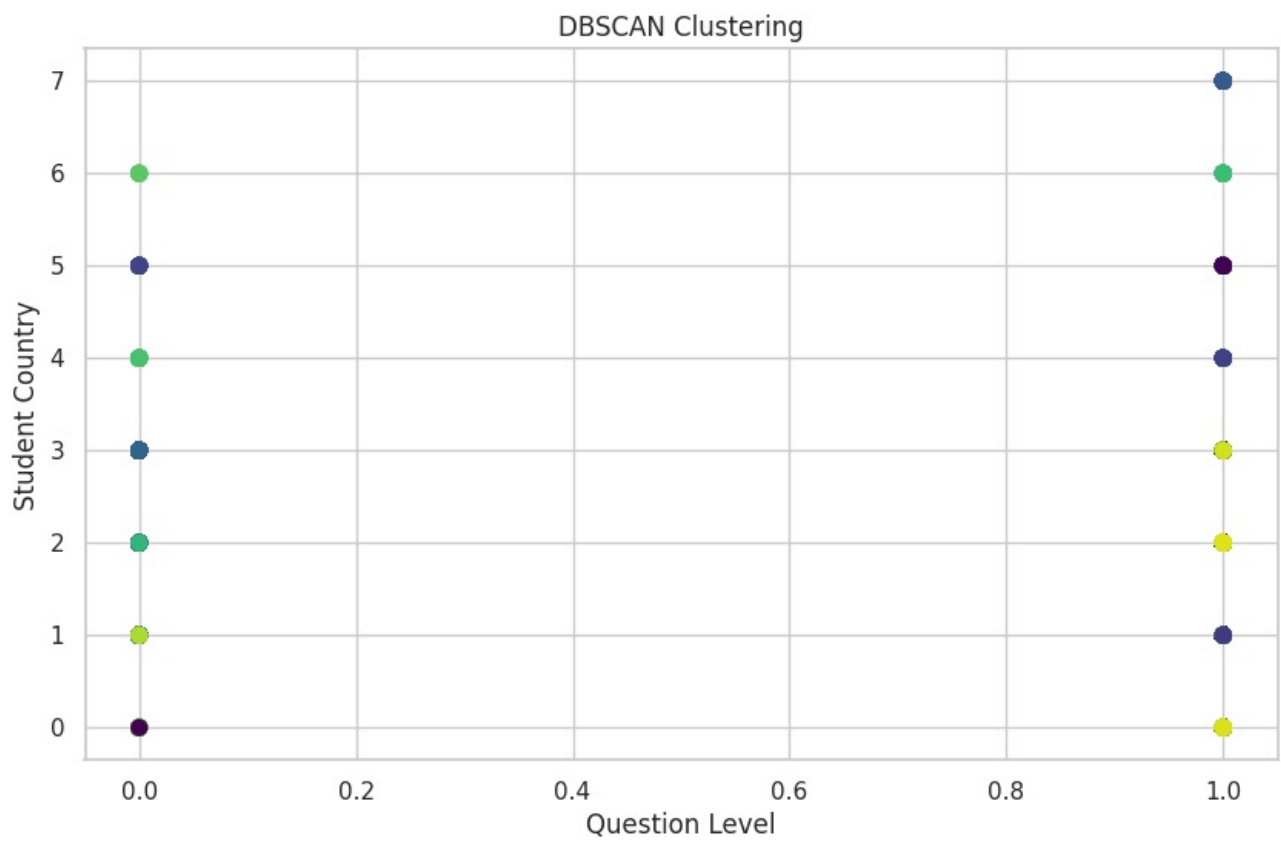
Cluster

0 0.455584 3940

1 0.428339 1842

2 0.438878 2495

Hierarchical Silhouette Score: 0.6052

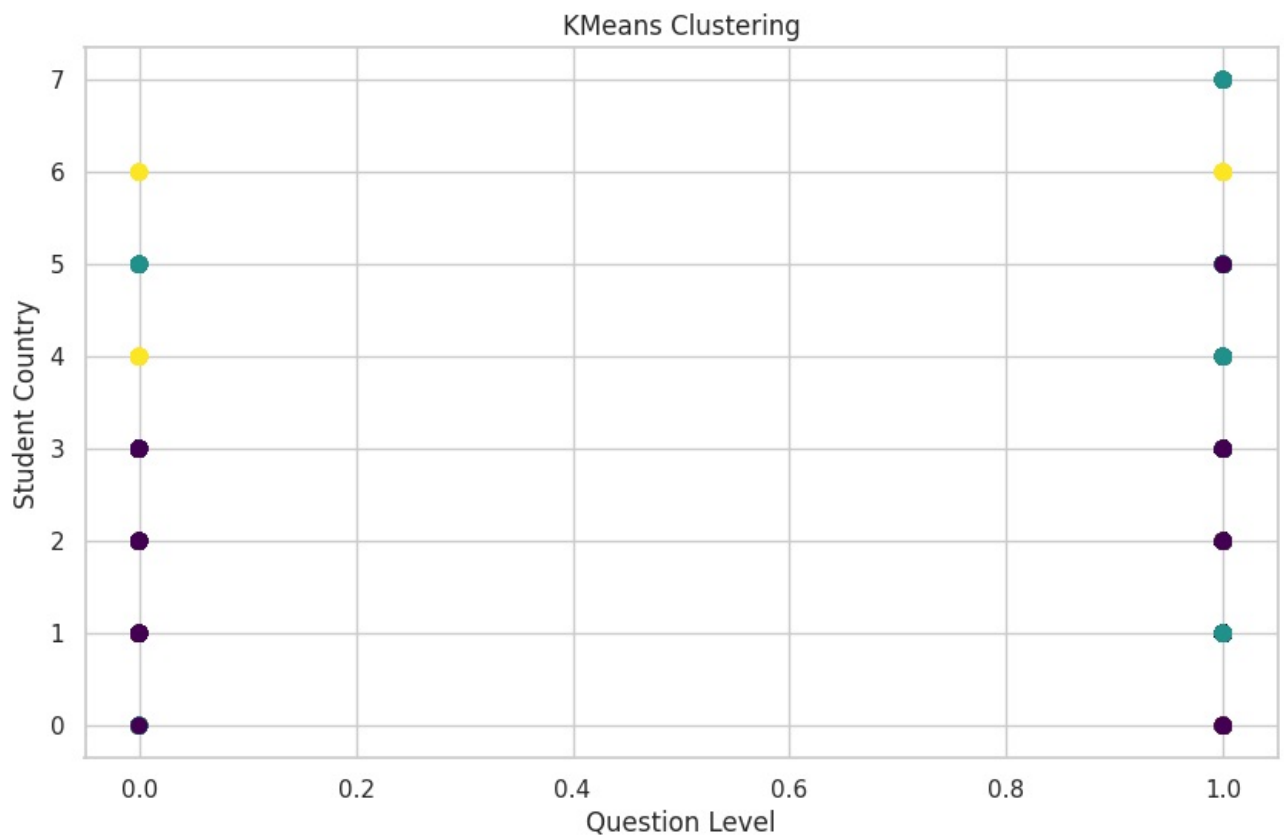


DBSCAN Cluster Analysis:

Type of Answer		
Cluster	mean	count
-1	0.375000	48
0	0.402985	67
1	0.531100	209
2	0.339623	53
3	0.500000	8
...
100	0.510638	47
101	0.523810	21
102	0.857143	7
103	0.714286	14
104	0.733333	15

[106 rows x 2 columns]

DBSCAN Silhouette Score: 0.8970



KMeans Cluster Analysis:

Type of Answer	mean	count
----------------	------	-------

Cluster

0	0.457828	3794
---	----------	------

1	0.429570	1860
---	----------	------

2	0.435761	2623
---	----------	------

KMeans Silhouette Score: 0.6170

Summary of Silhouette Scores

Hierarchical Clustering: 0.6051609243237689

DBSCAN: 0.8969959437931226

KMeans: 0.6170315937621726

In []: